

LLaMAR: Long-Horizon Planning for Multi-Agent Robotics in Partially Observable Environments


Siddharth Nayak*, Adelmo Orozco*, Marina Have, Vittal Thirumalai, Jackson Zhang, Darren Chen, Aditya Kapoor, Eric Robinson, Karthik Gopalakrishnan, James Harrison, Brian Ichter, Anuj Mahajan, Hamsa Balakrishnan




Motivation



Motivation




Motivation




Motivation




Motivation



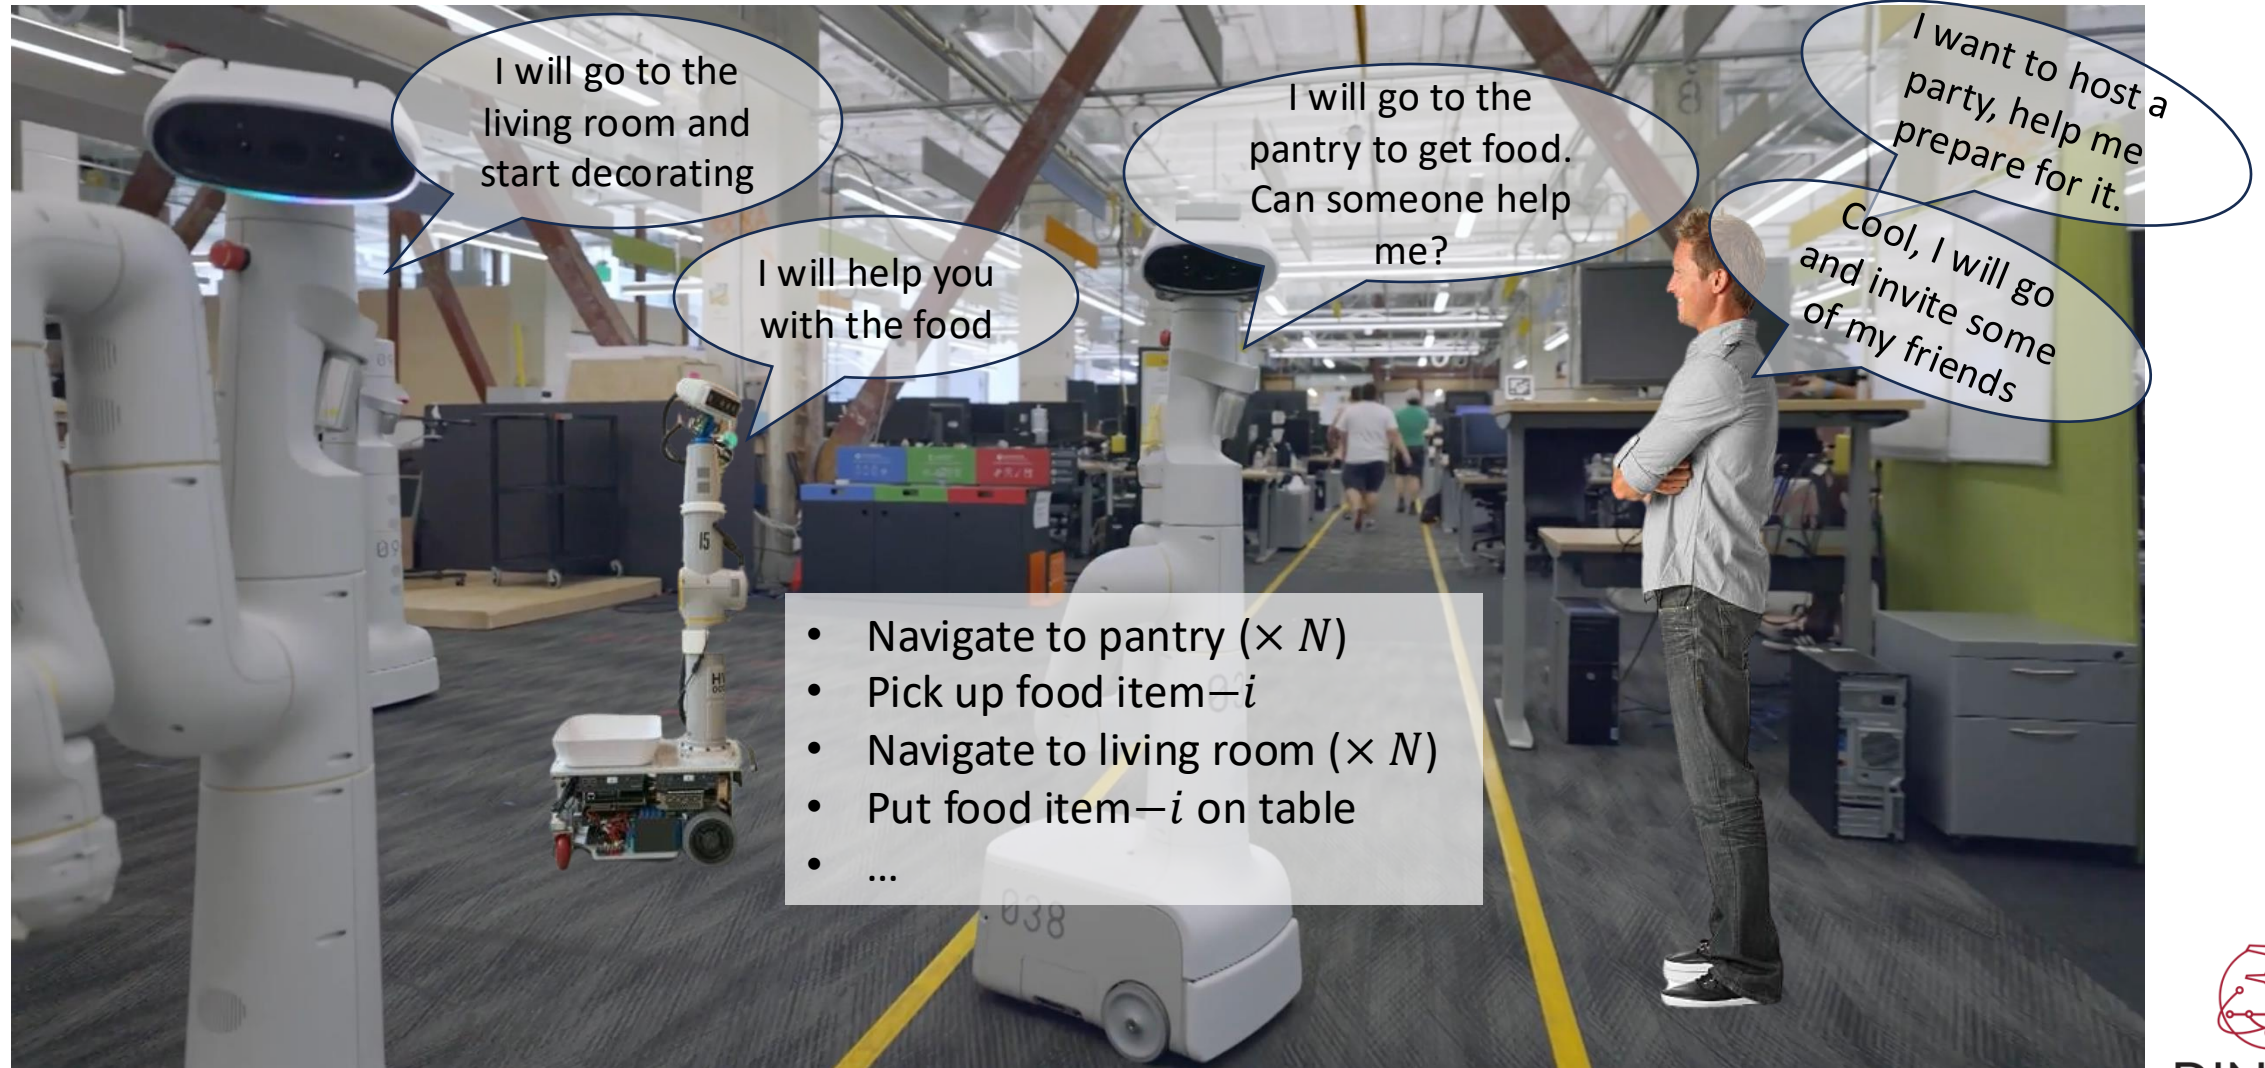
Motivation



Long-Horizon Planning



Long-Horizon Planning



Long-Horizon Planning

I will go to the living room and start decorating

I will help you with the food

I will go to the pantry to get food. Can someone help me?

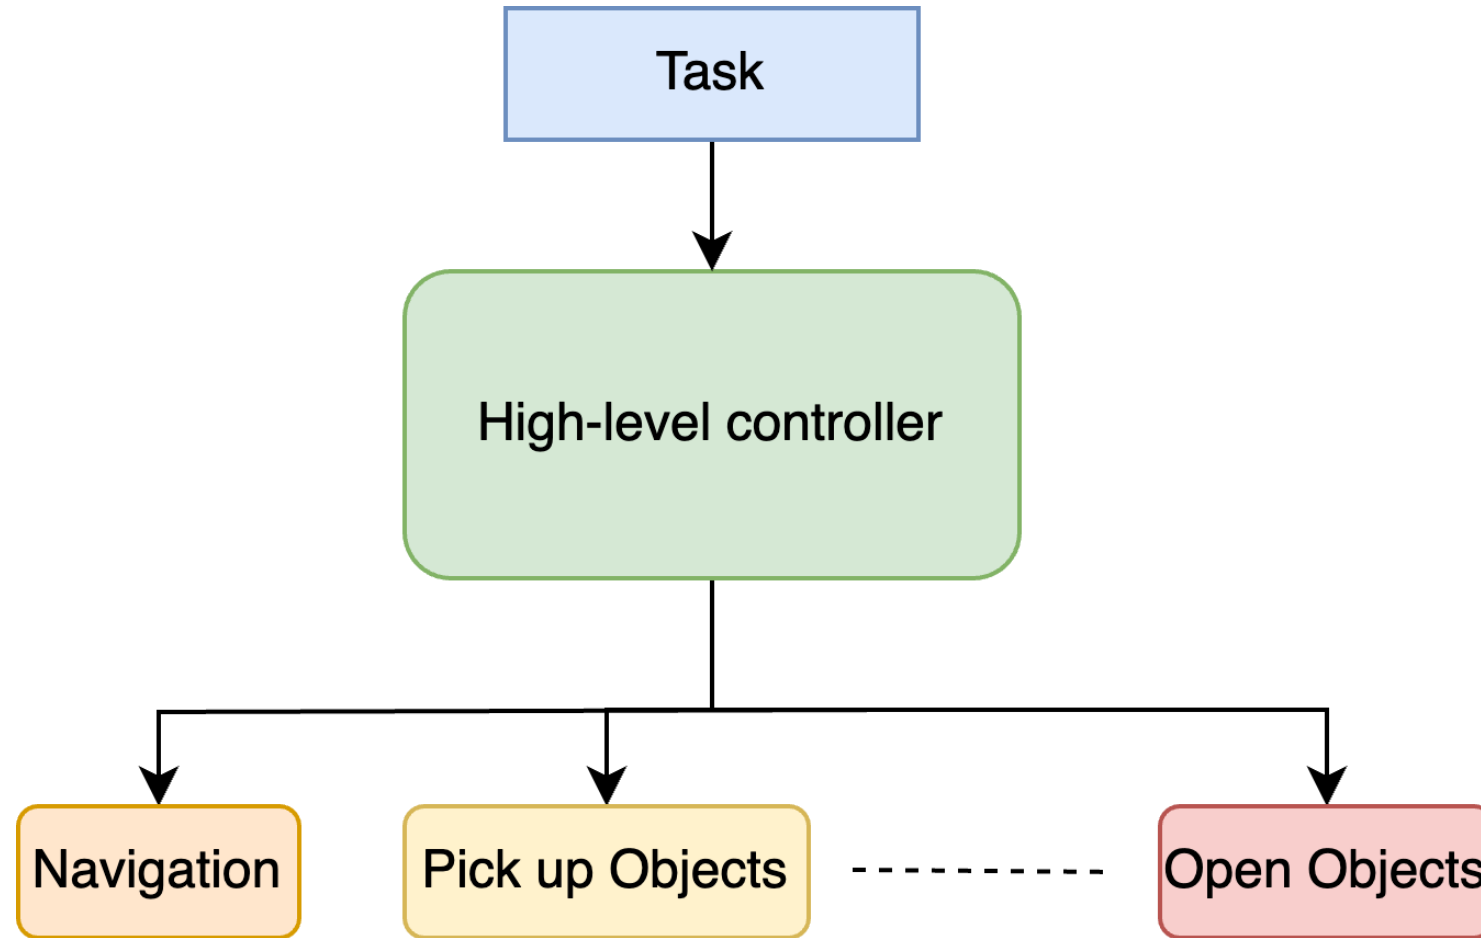
Cool, I will go and invite some of my friends

I want to host a party, help me prepare for it.

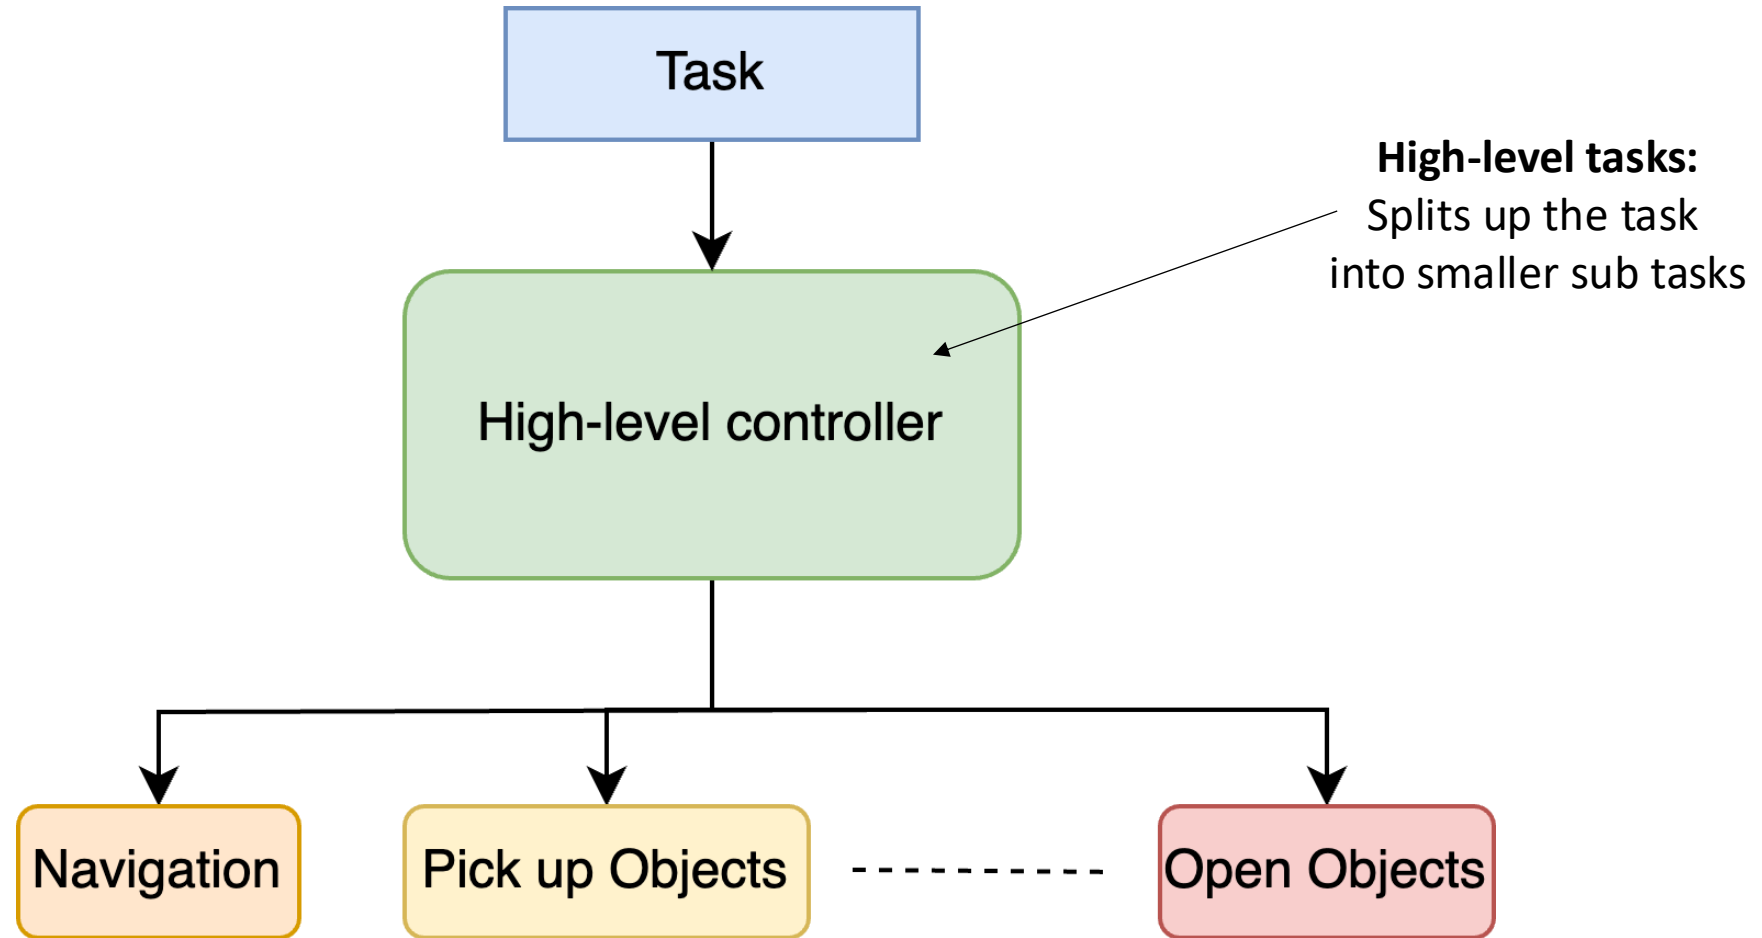
- Navigate to living room
- Navigate to lightSwitch— i
- Switch on light— i
- Navigate to basement ($\times N$)
- Pick up decoration item — i
- ...

- Navigate to pantry ($\times N$)
- Pick up food item— i
- Navigate to living room ($\times N$)
- Put food item— i on table
- ...

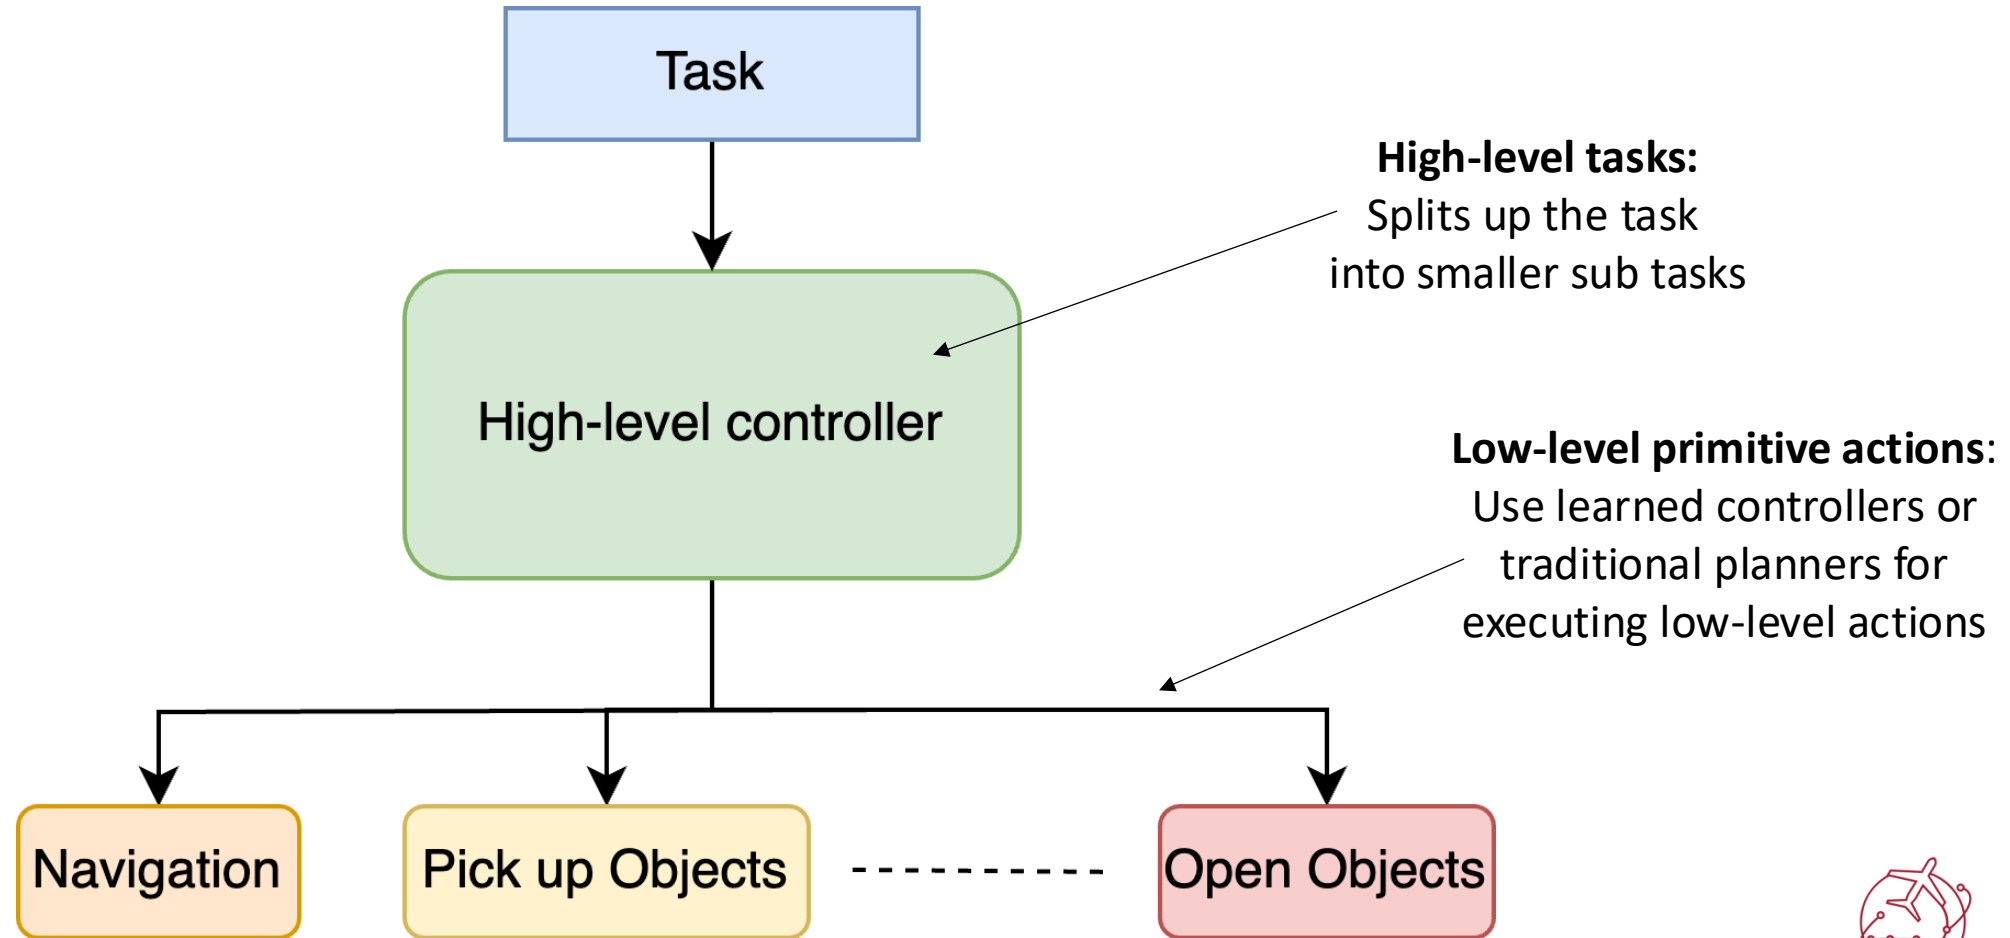
Hierarchical Planning



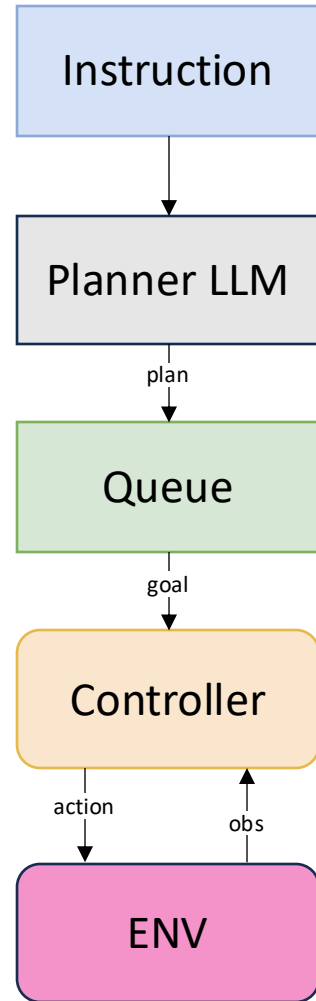
Hierarchical Planning




Hierarchical Planning



LLMs for planning




LLMs for planning



- Does not accommodate for uncertainties



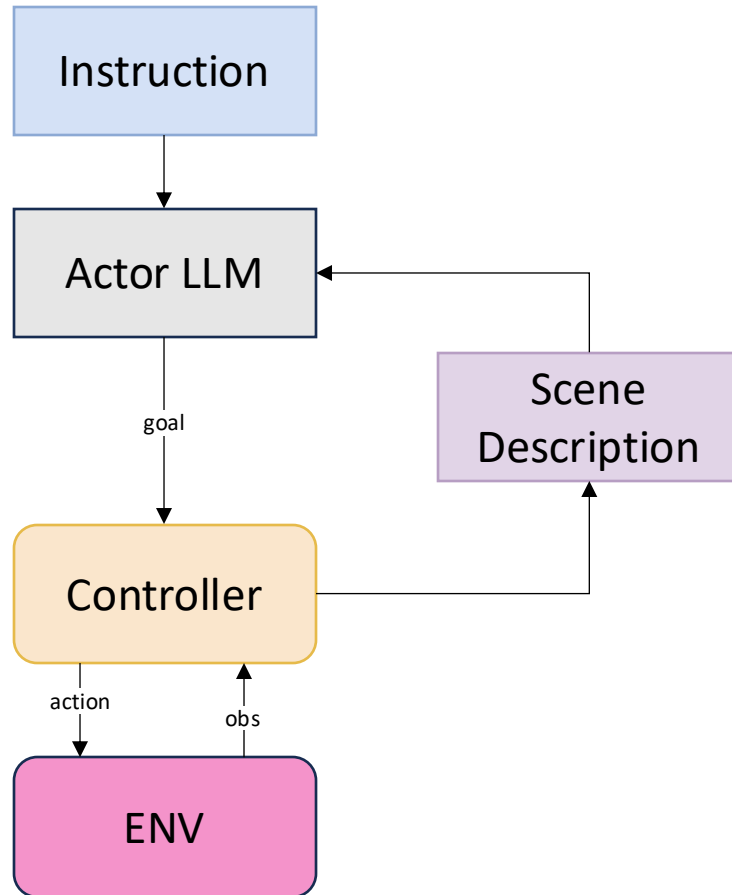
LLMs for planning



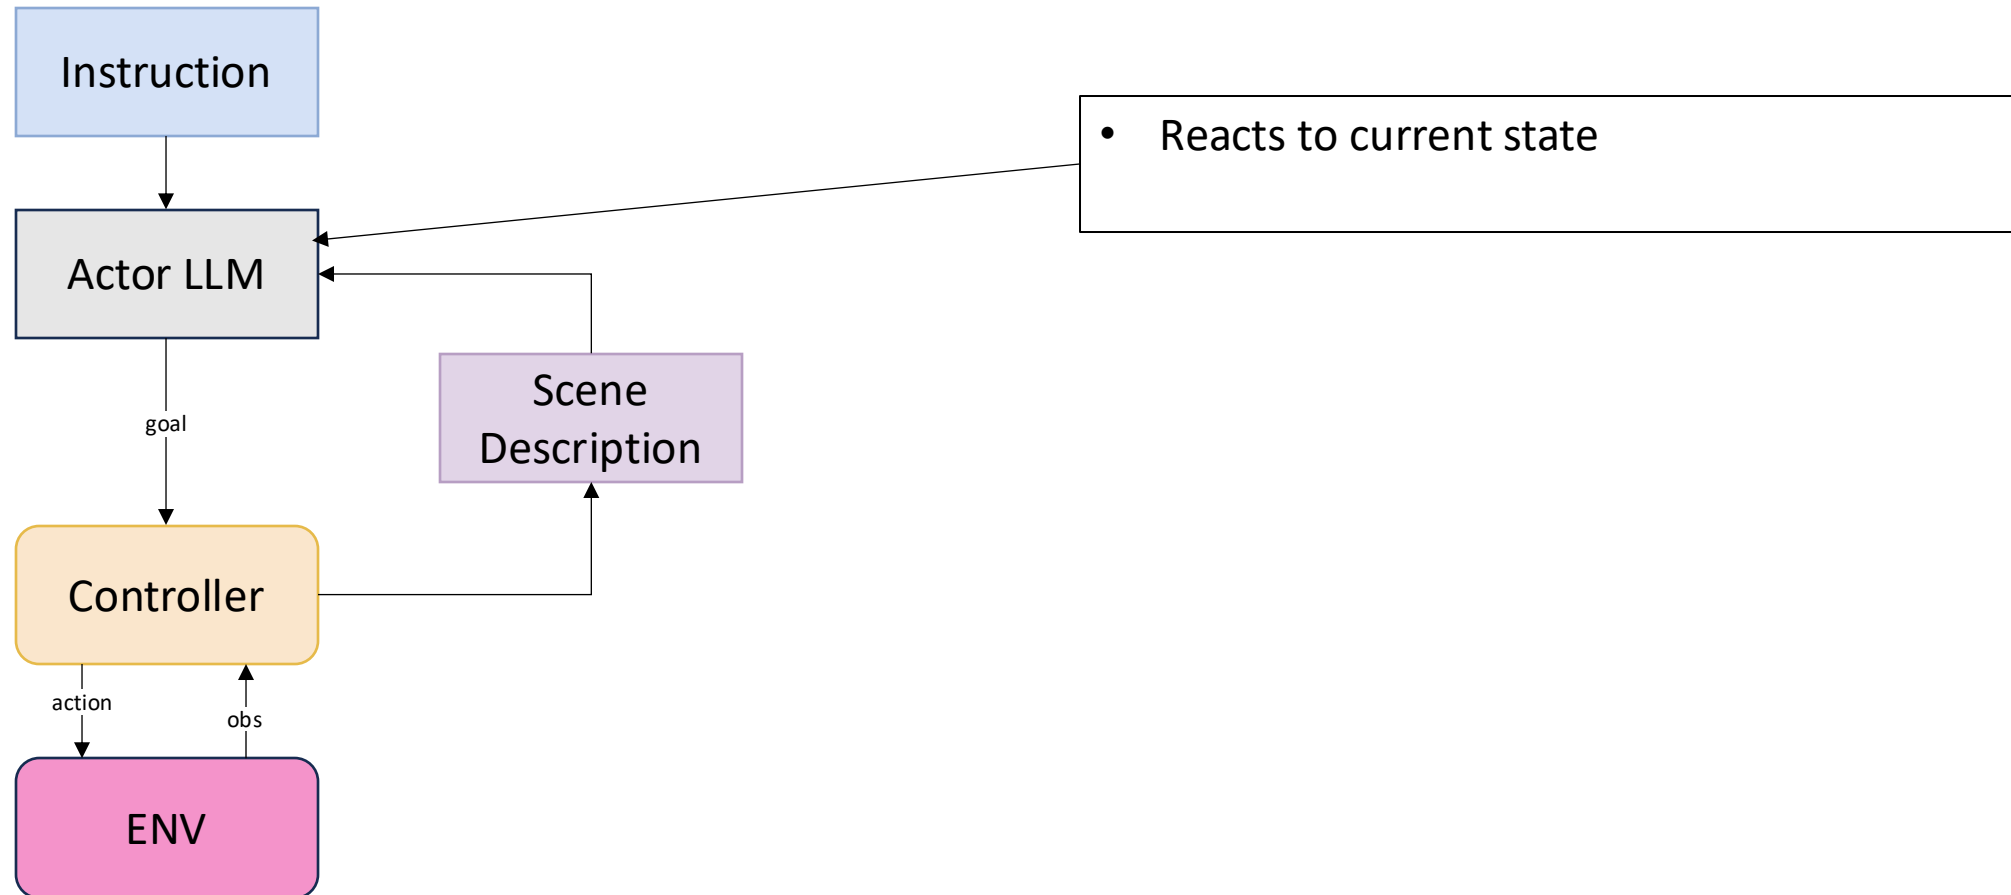
- Does not accommodate for uncertainties
- Requires prior knowledge about subtasks (global information)



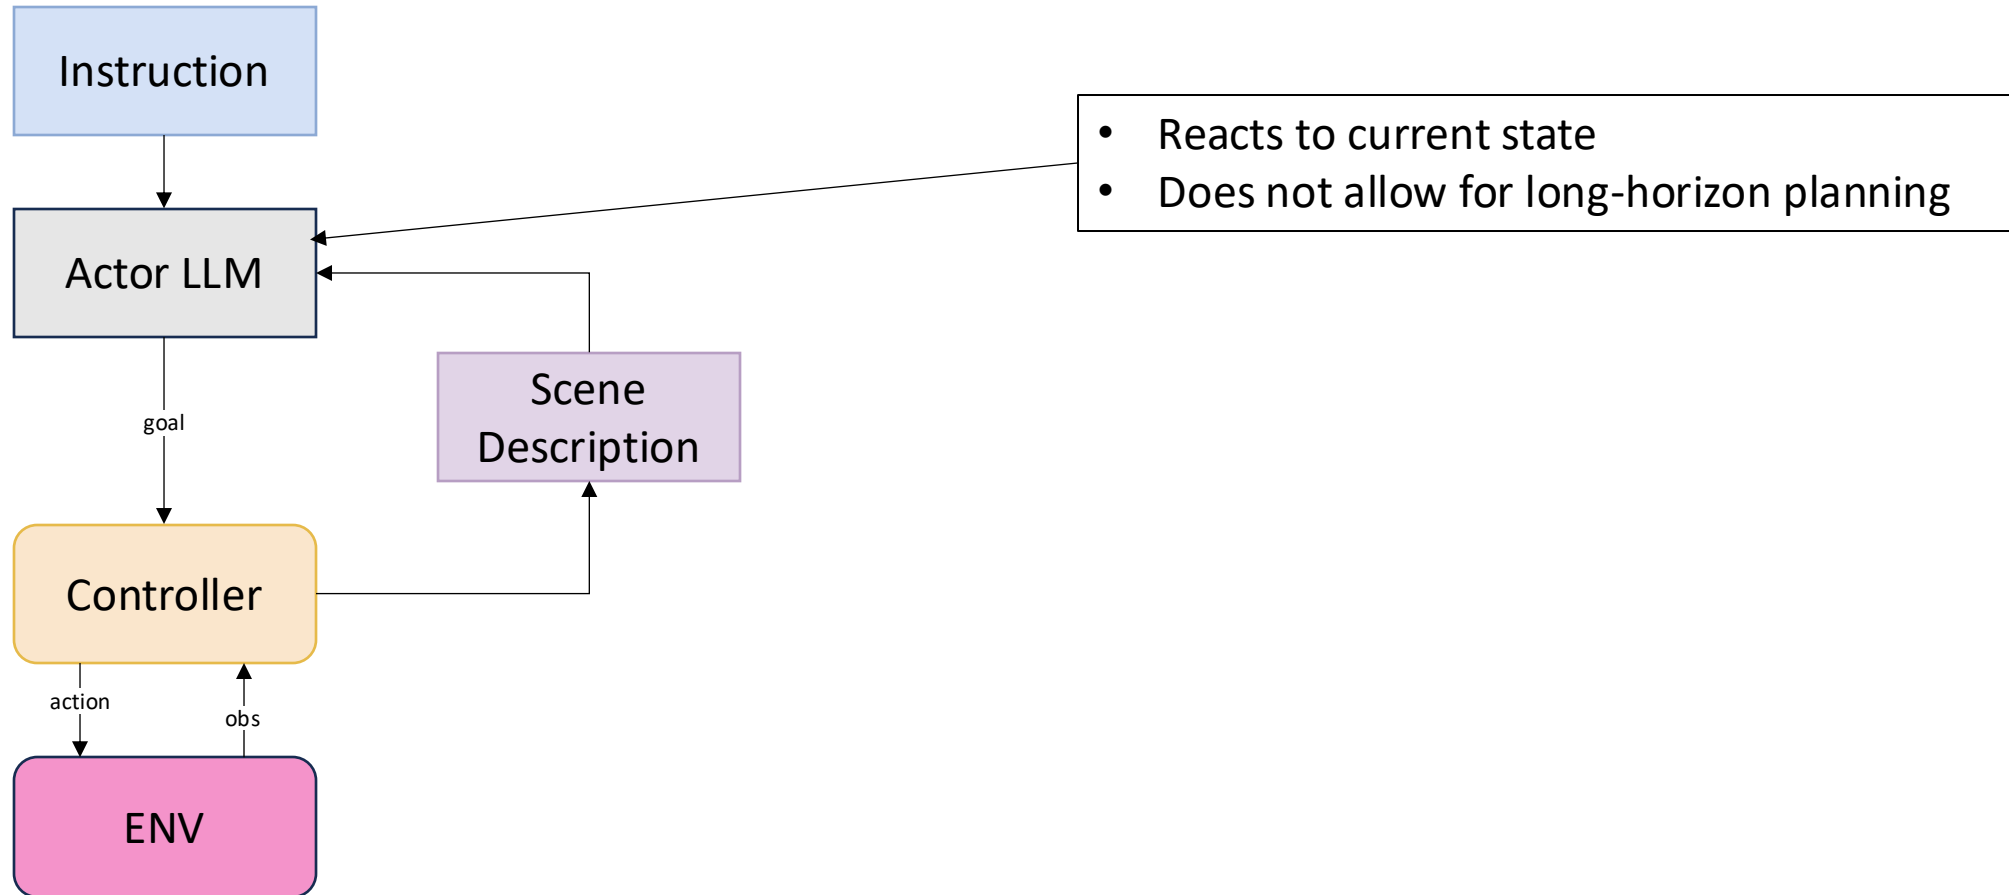
LLMs for Reactive Planning



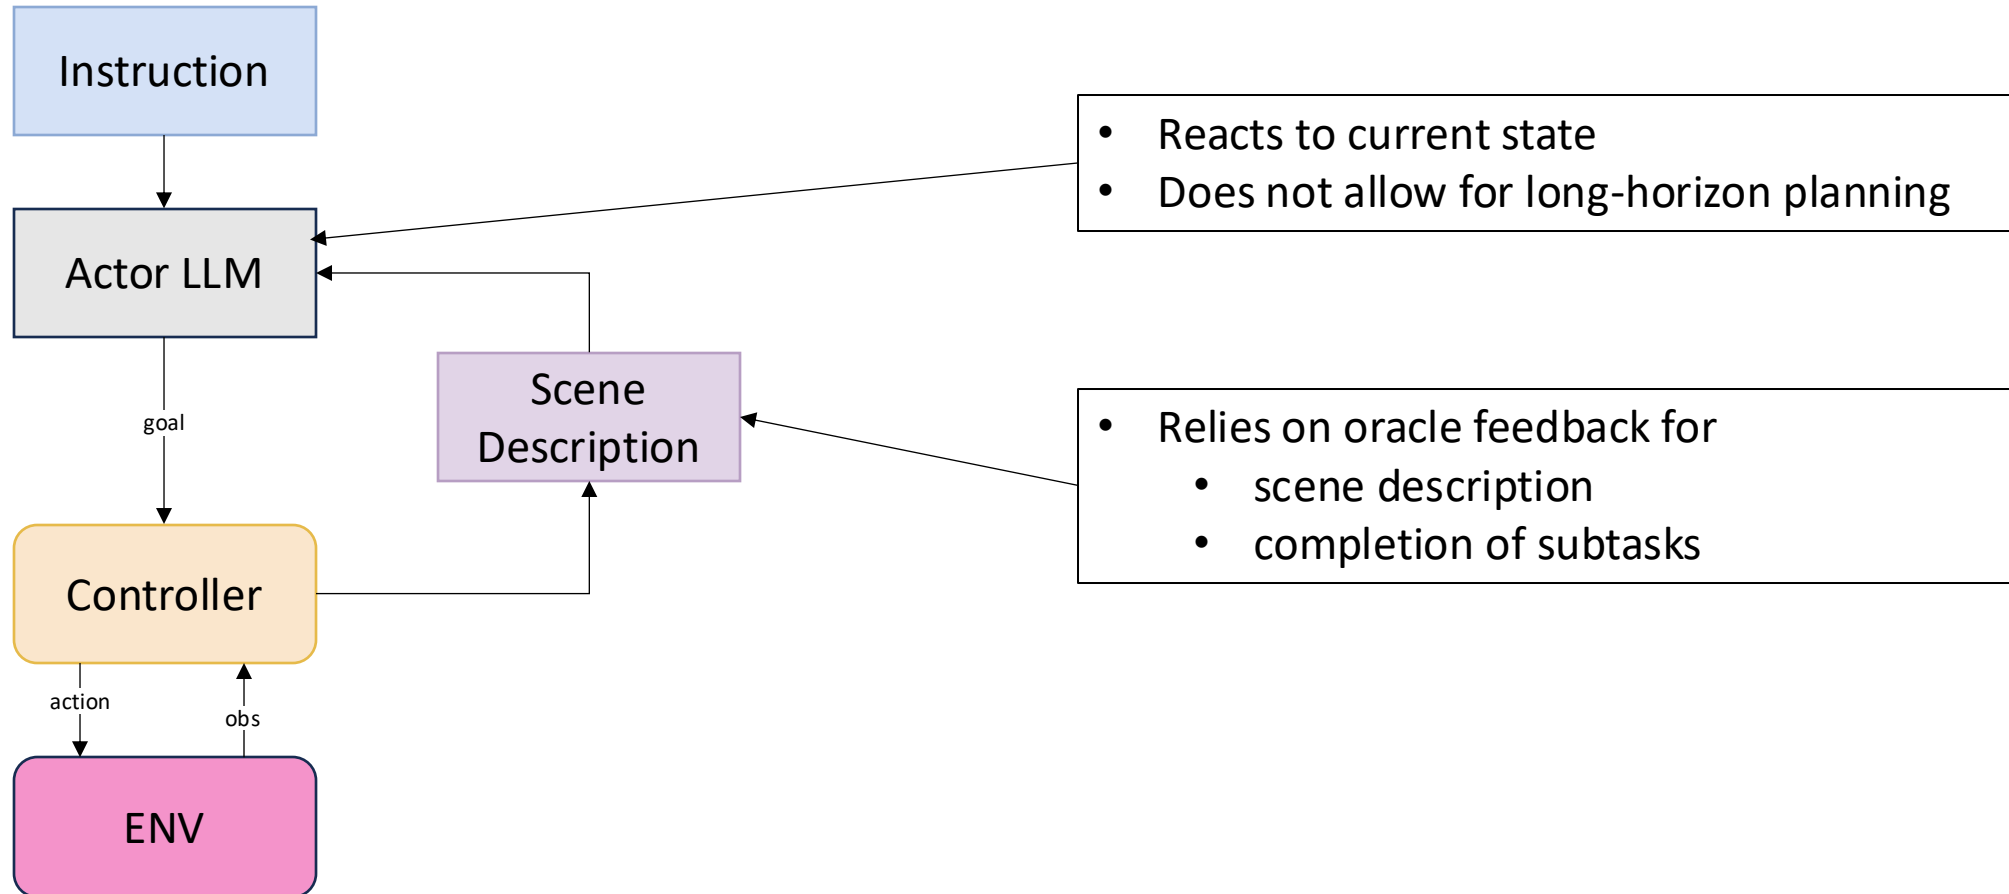
LLMs for Reactive Planning



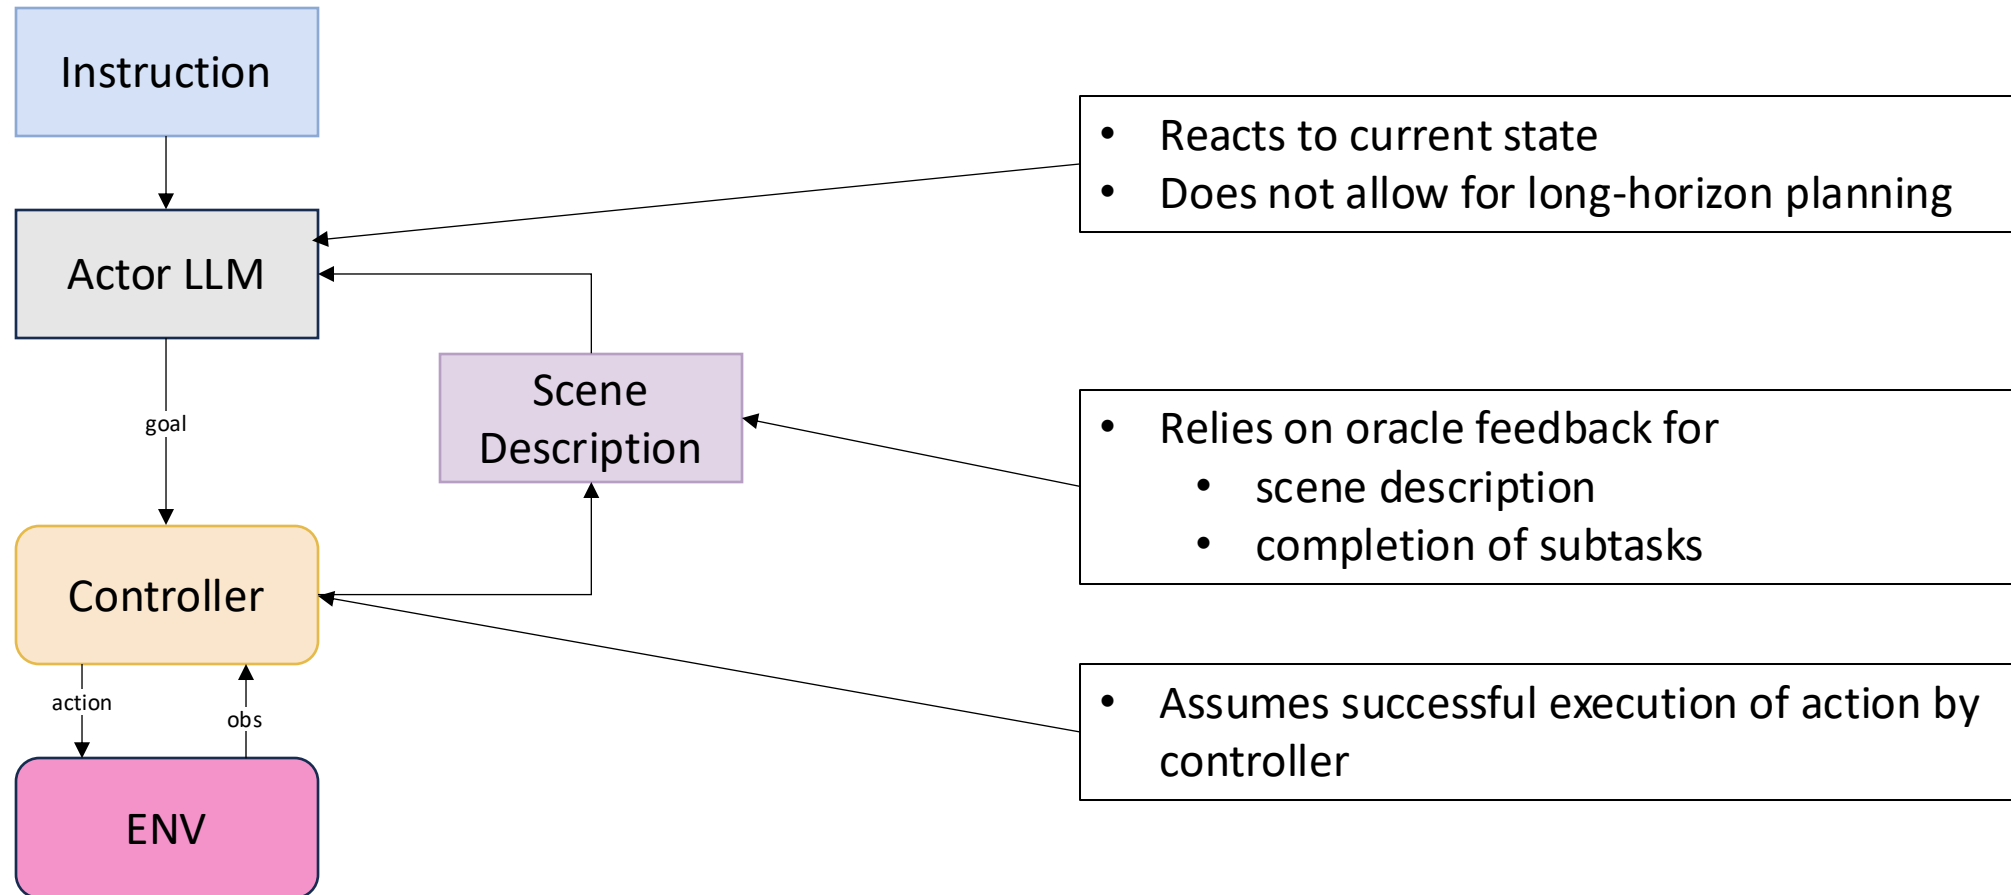
LLMs for Reactive Planning



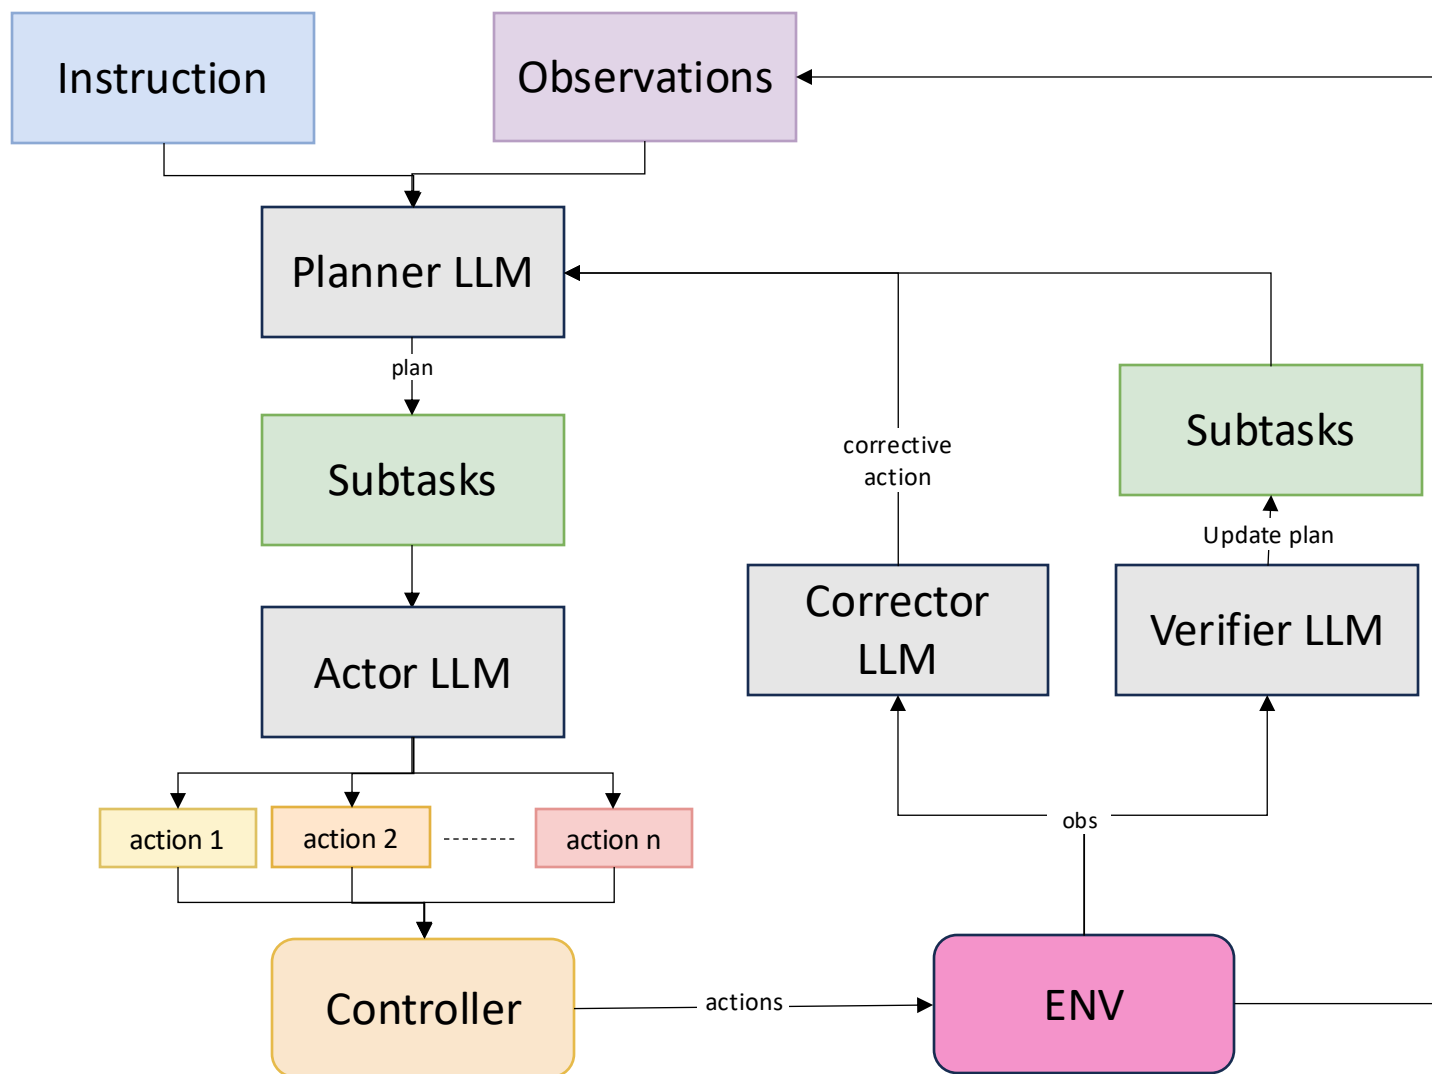
LLMs for Reactive Planning



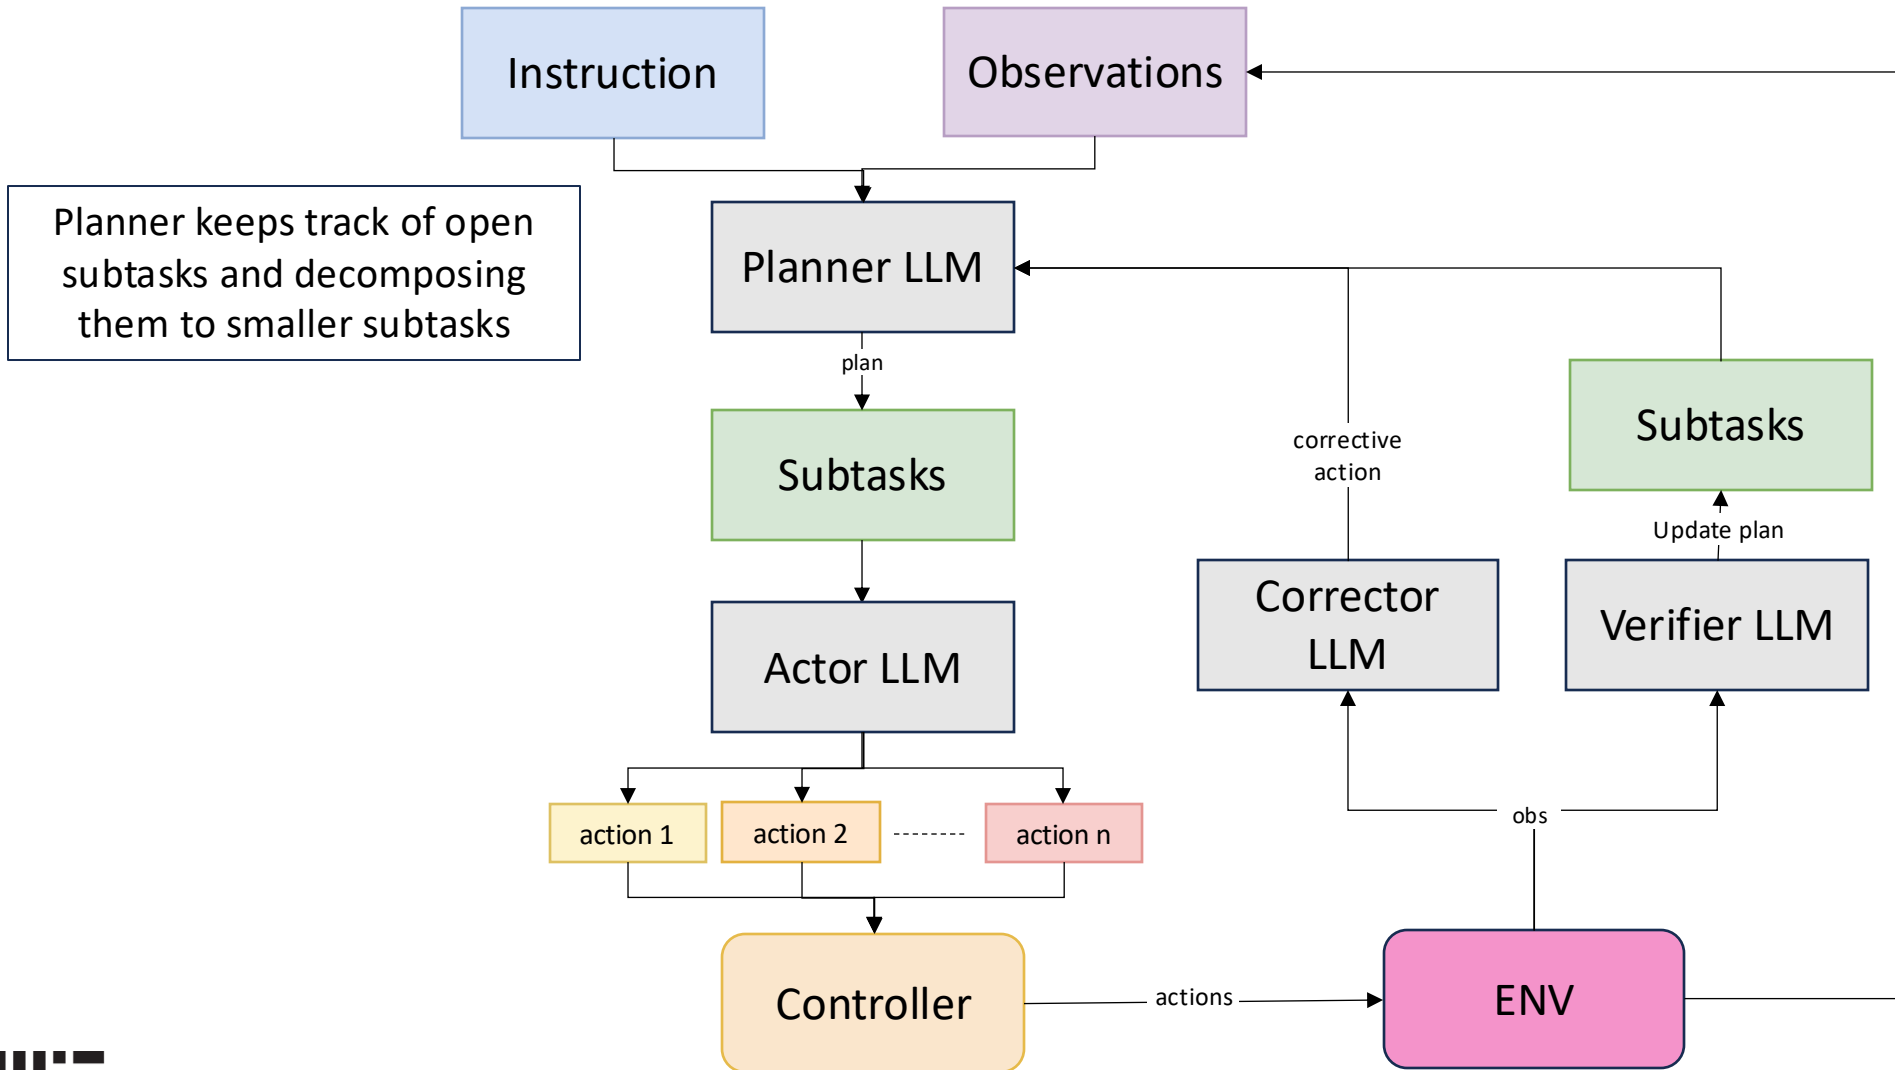
LLMs for Reactive Planning



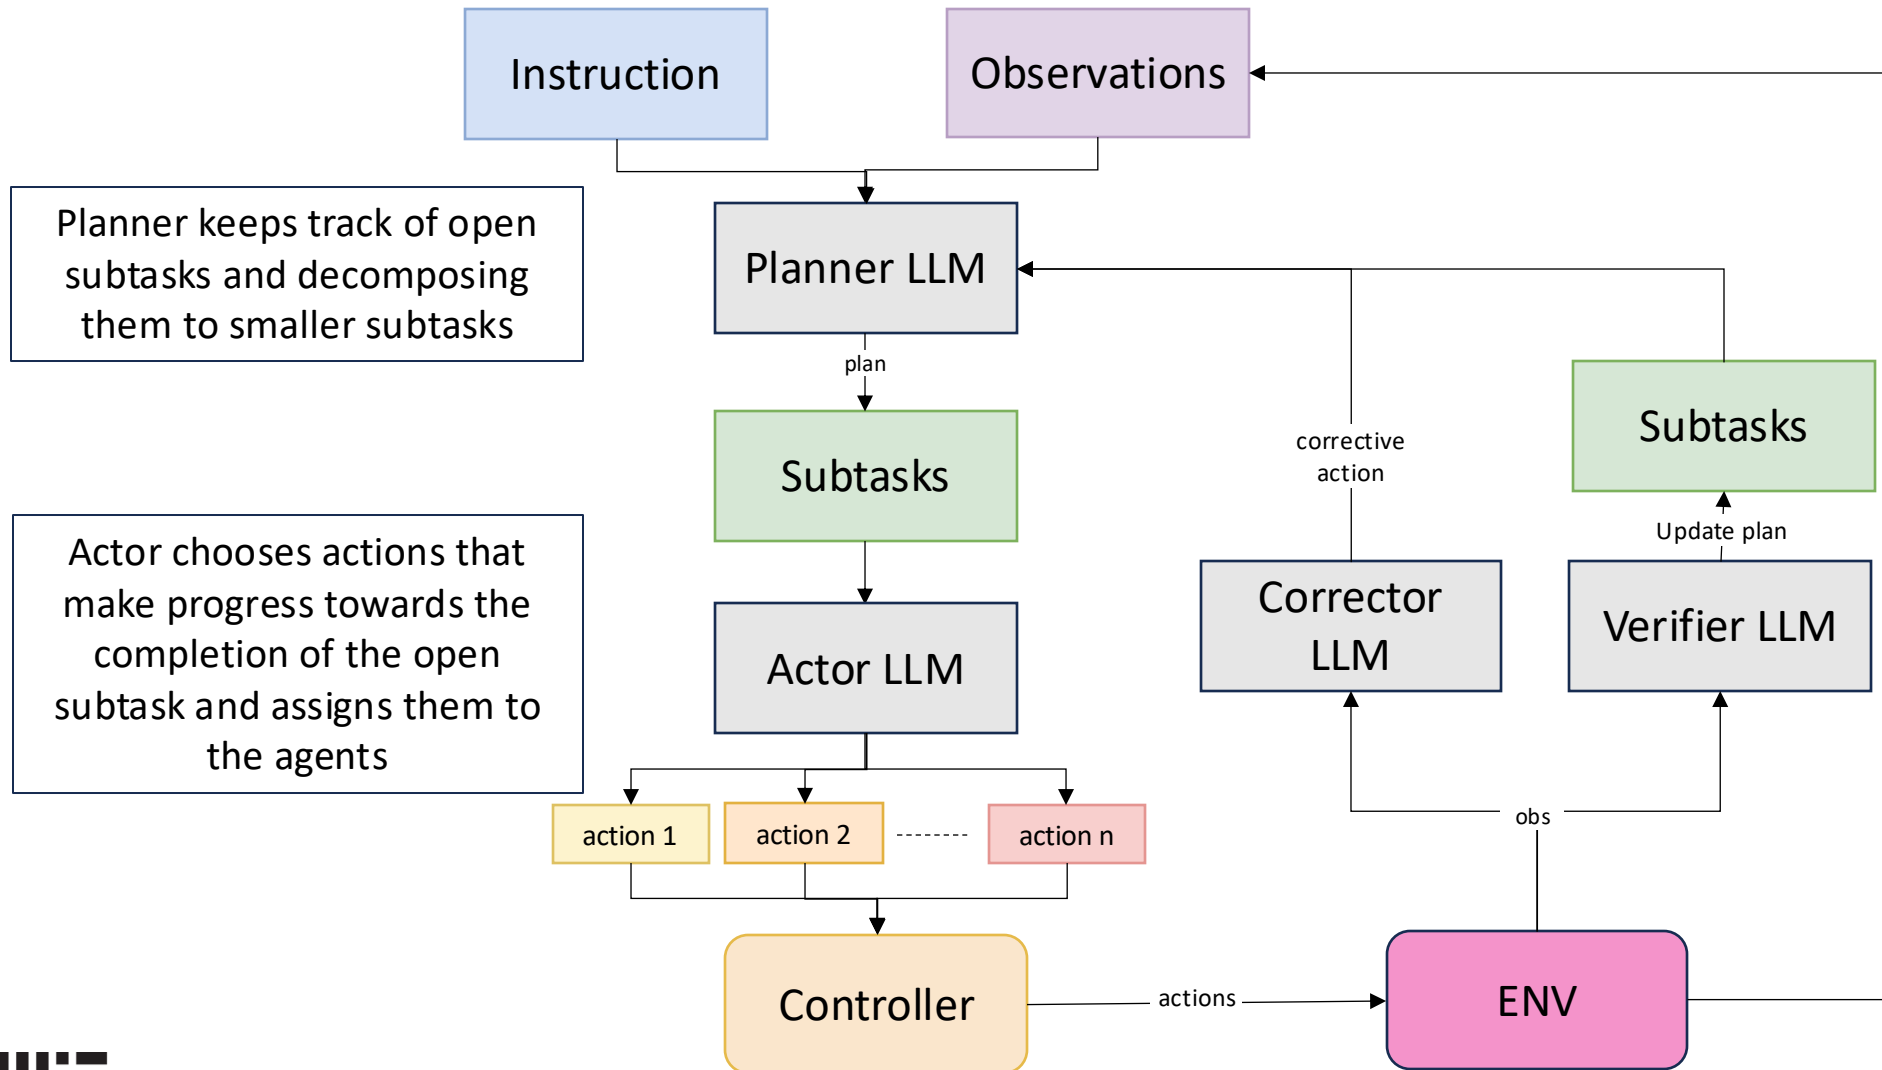
LLaMAR



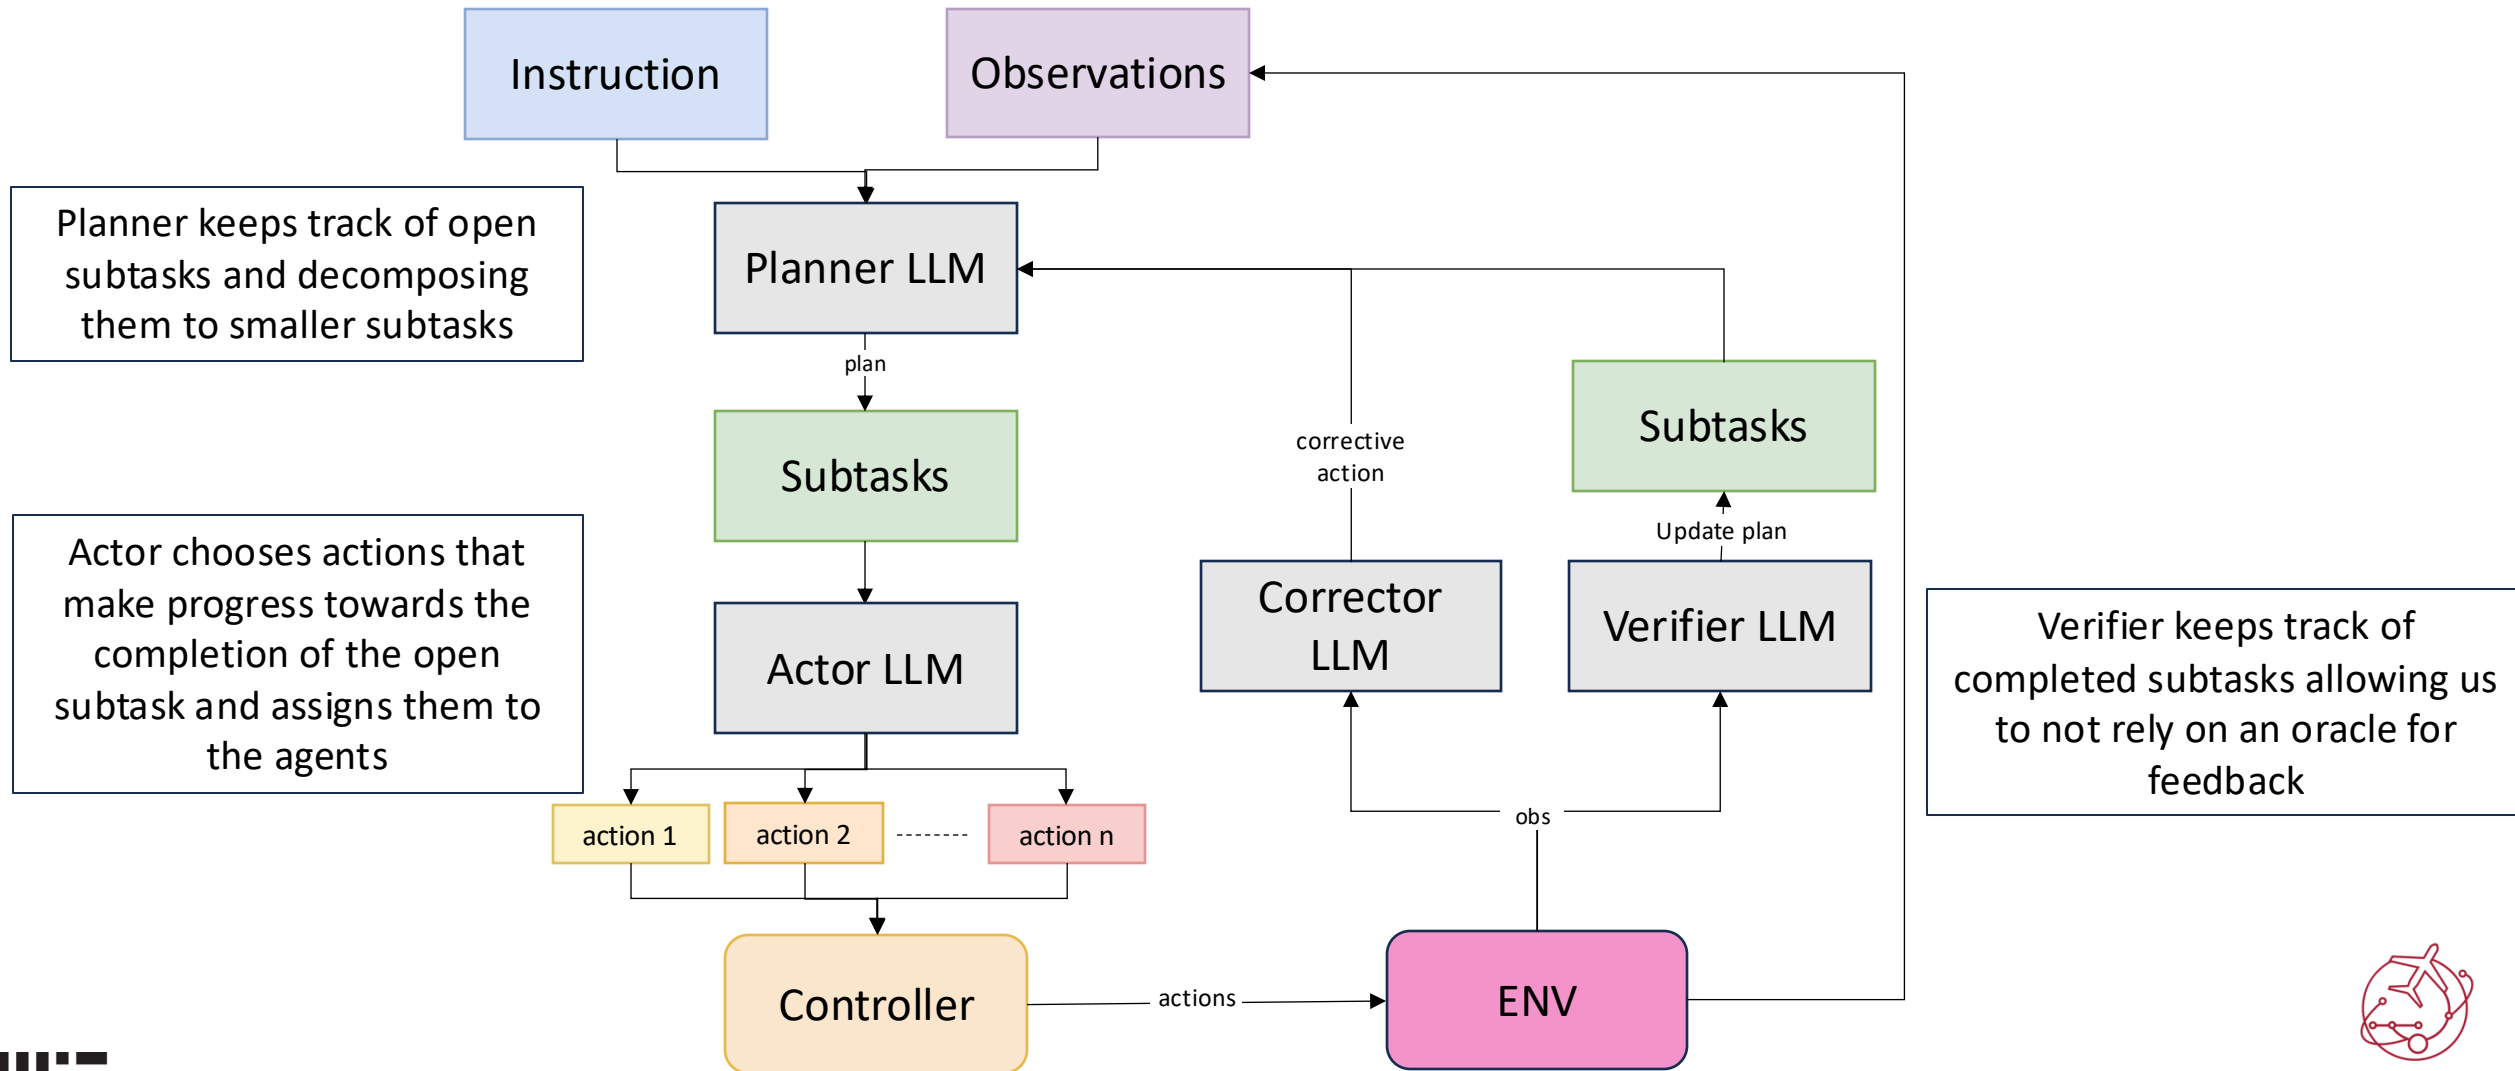
LLaMAR



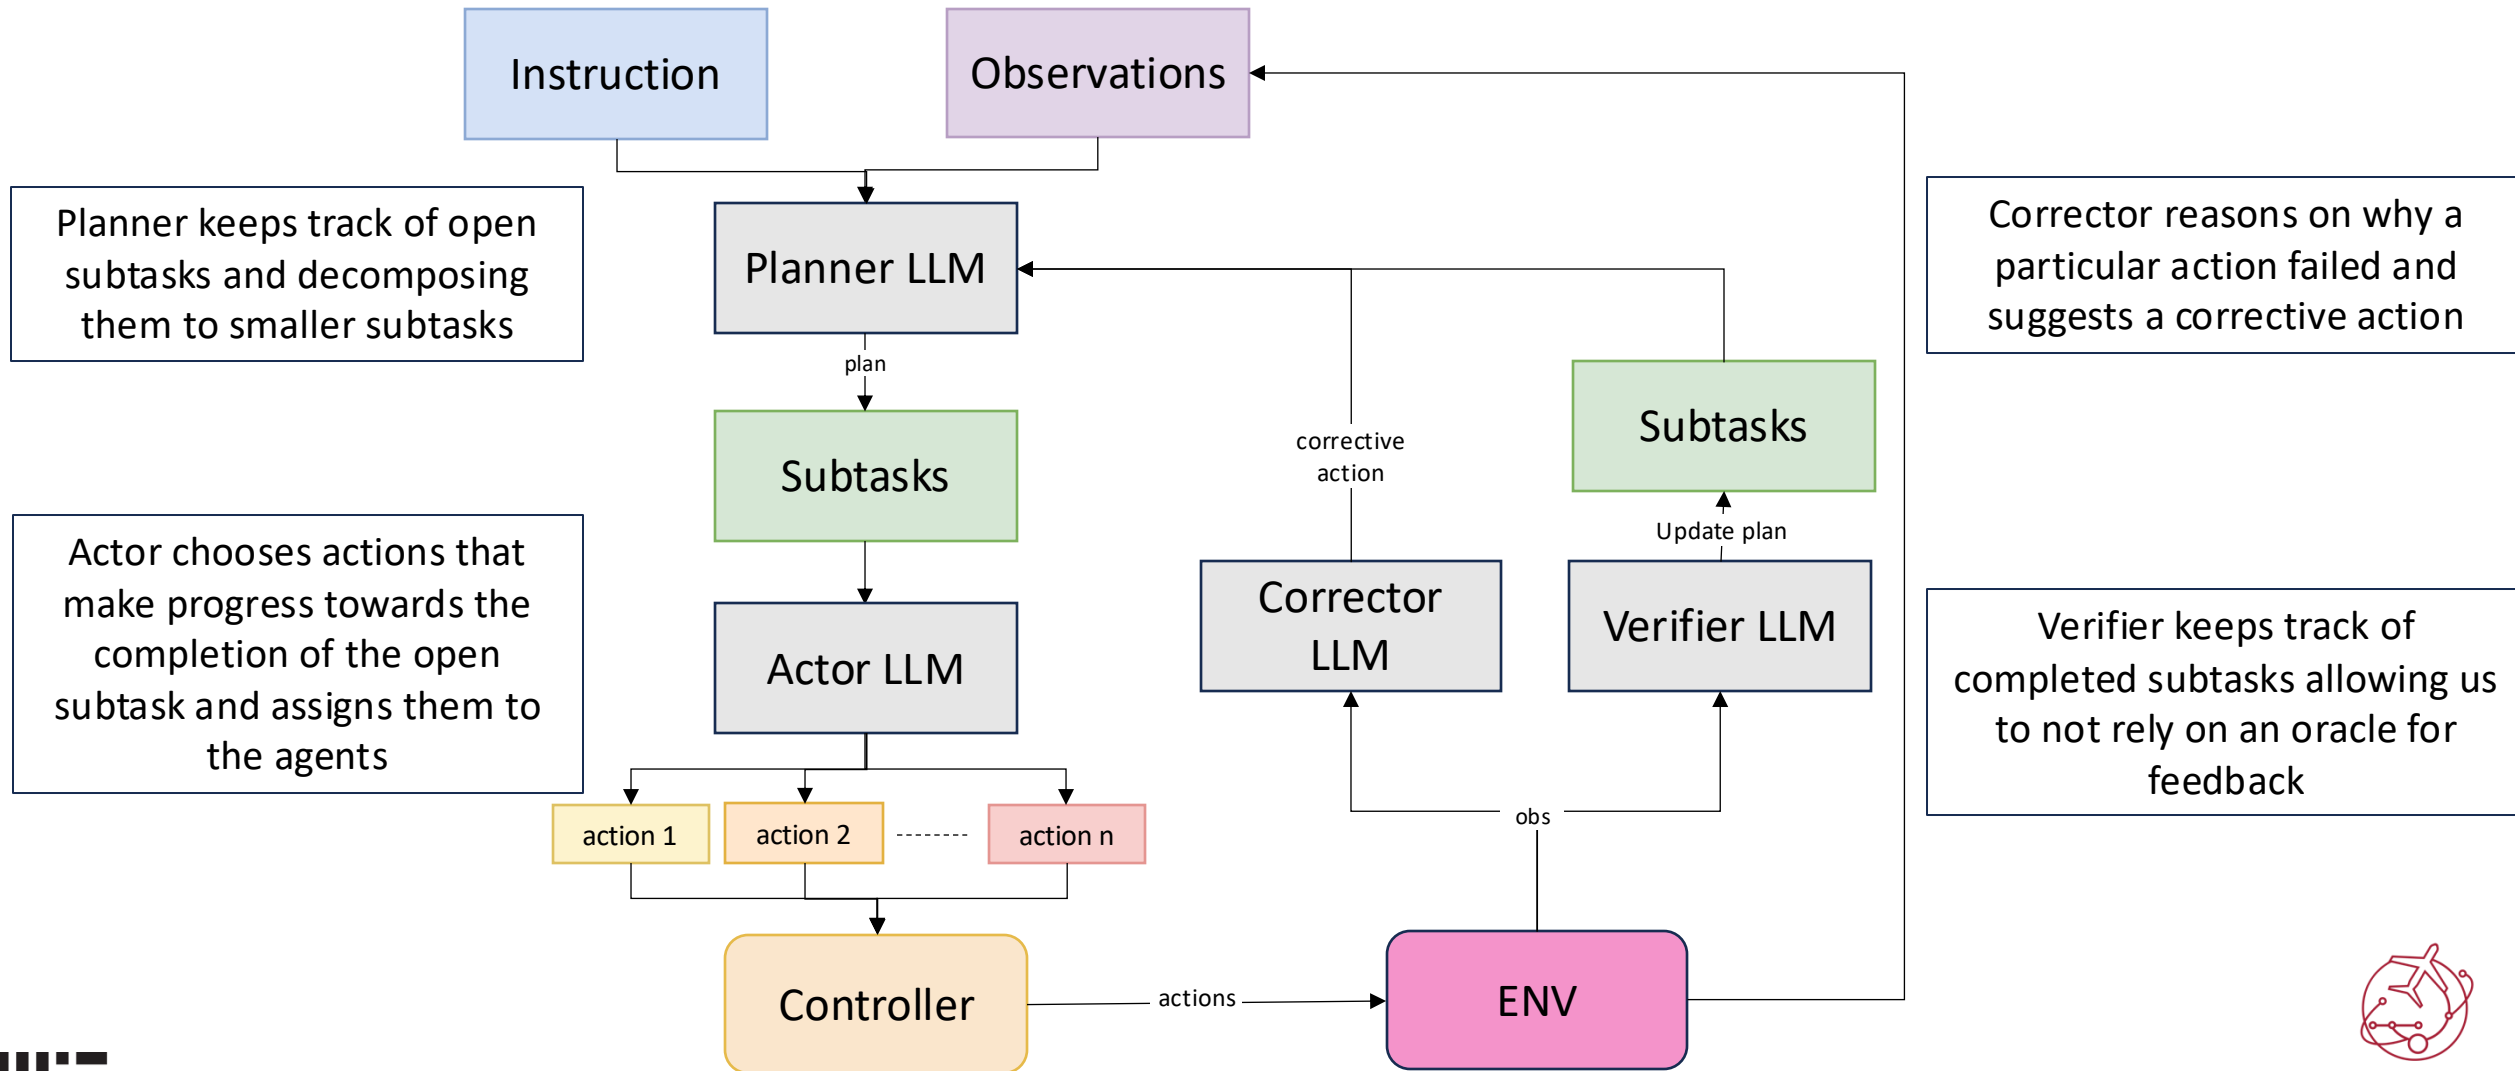
LLaMAR



LLaMAR



LLaMAR



MAP-THOR



Tasks

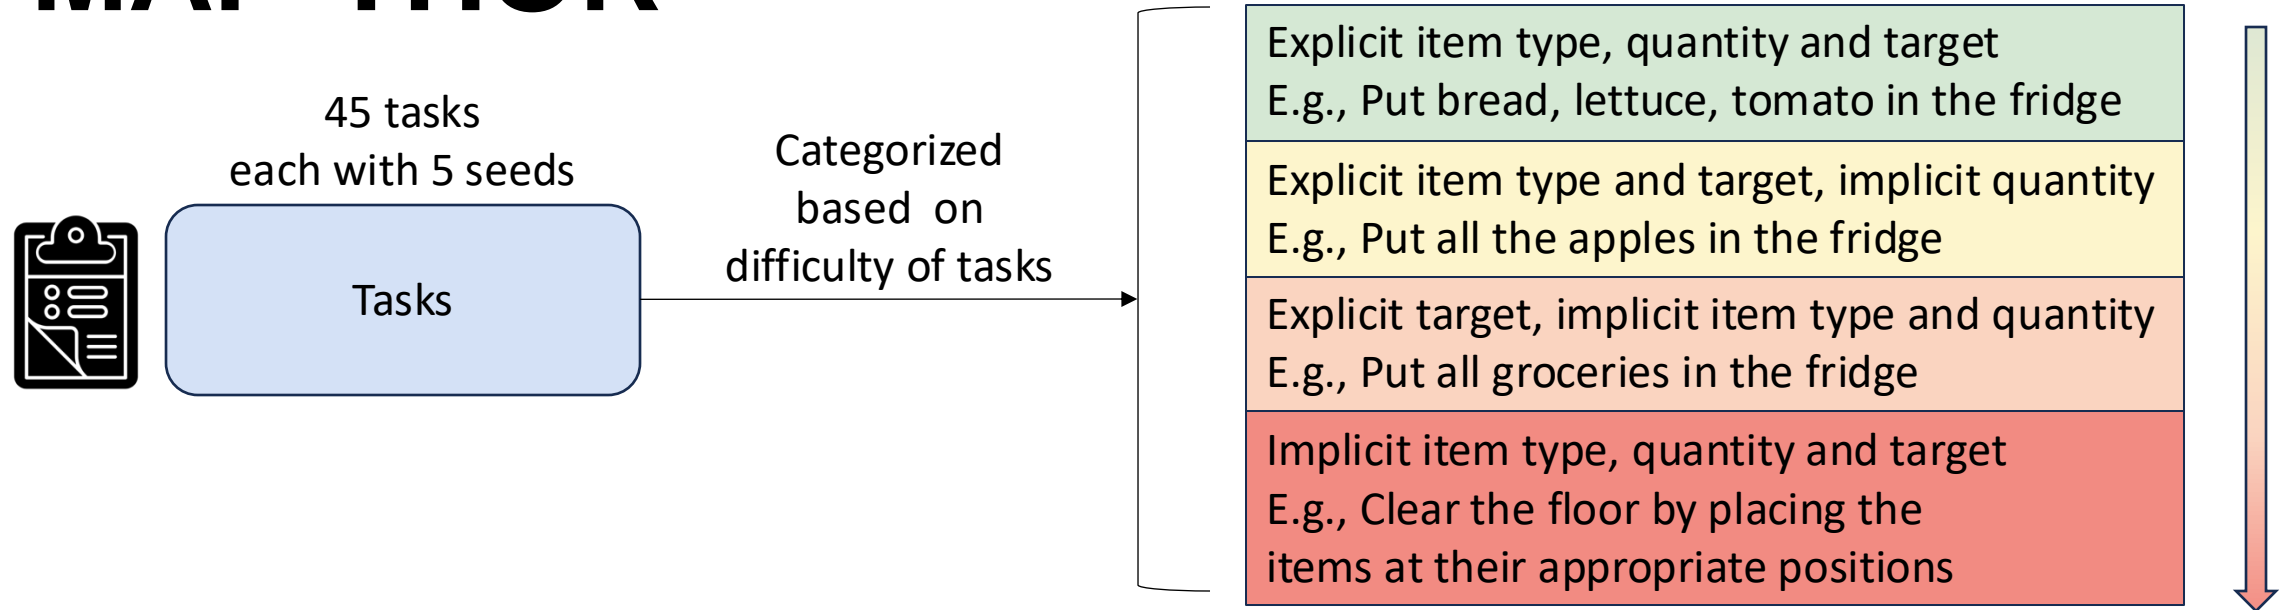
MAP-THOR

45 tasks
each with 5 seeds

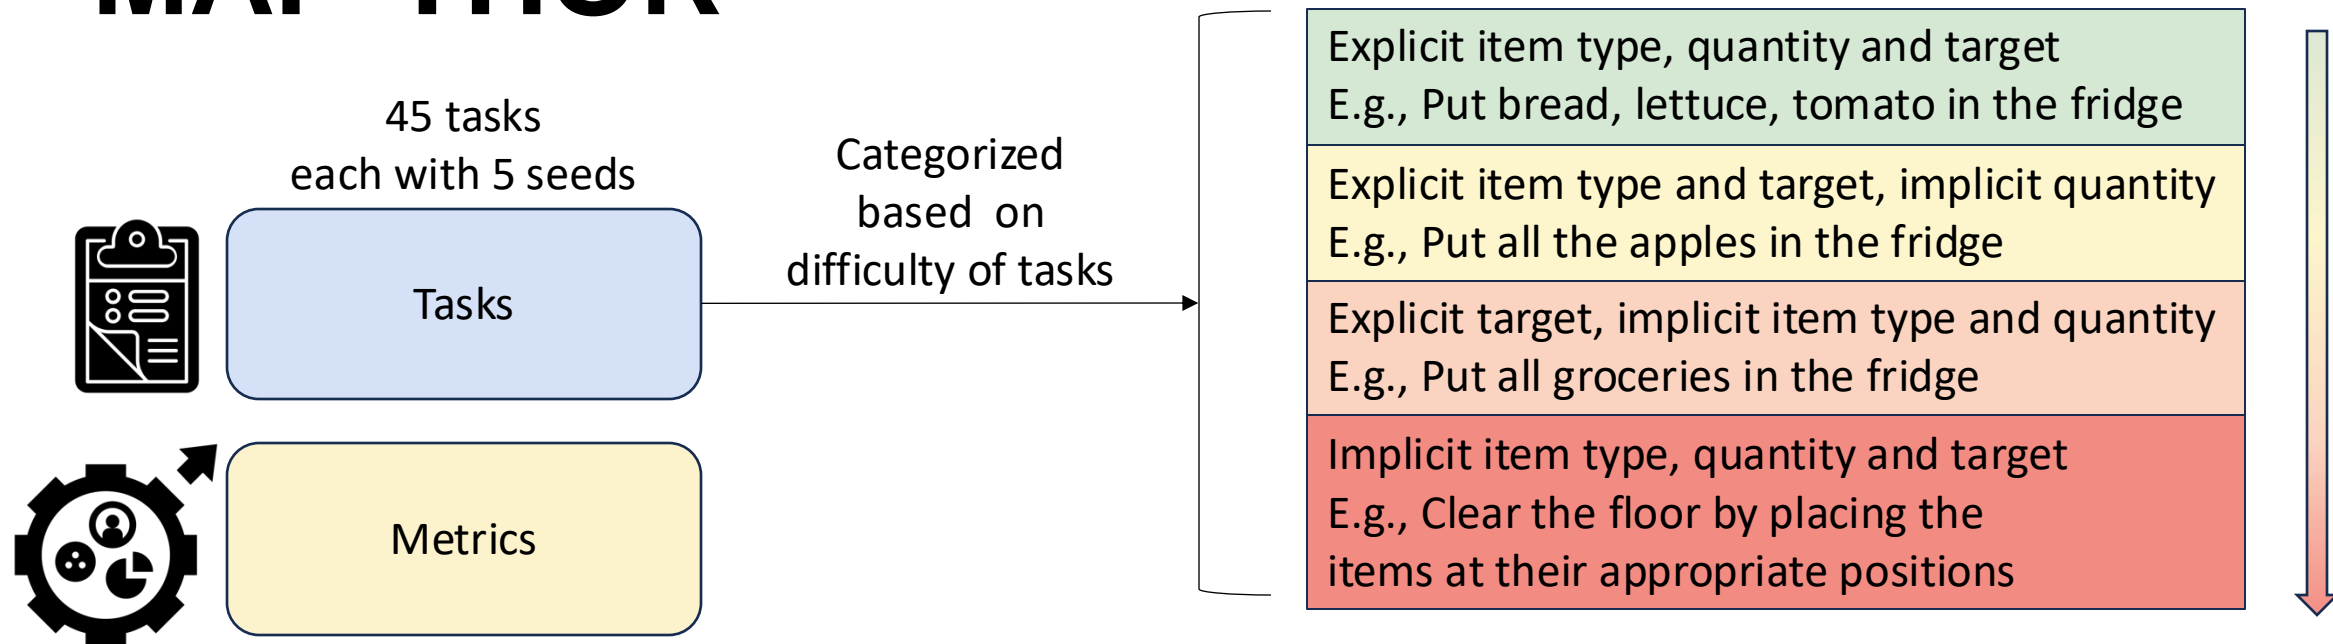


Tasks

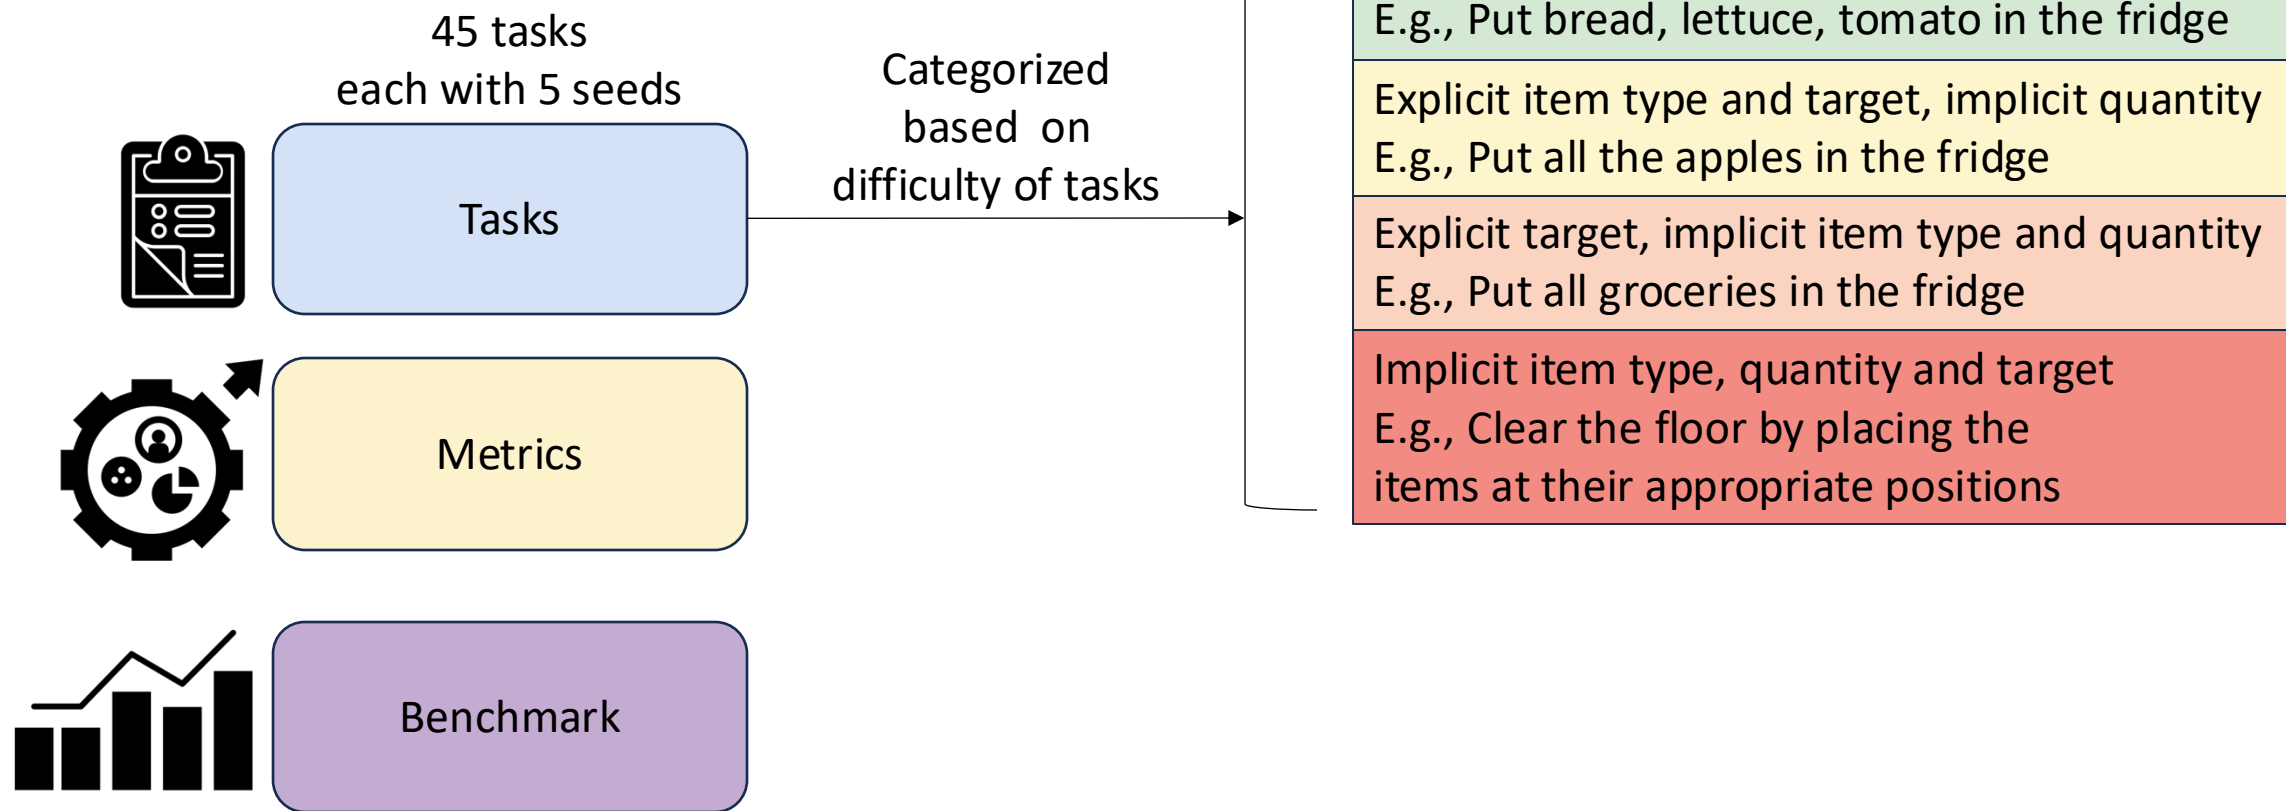
MAP-THOR



MAP-THOR



MAP-THOR



Demo

Human Instruction: "I want to put the groceries in the fridge"



Alice's POV



Bob's POV

Key Takeaways

Key Takeaways

- LLaMAR performs 30% better than other LLM-based methods for multi-agent long-horizon planning

Key Takeaways

- LLaMAR performs 30% better than other LLM-based methods for multi-agent long-horizon planning
- Vision Language Models perform better than pure text-based language models

Key Takeaways

- LLaMAR performs 30% better than other LLM-based methods for multi-agent long-horizon planning
- Vision Language Models perform better than pure text-based language models
- A modular cognitive architecture with distinct roles helps with solving complex tasks.

Key Contributions:

Key Contributions:

- A **modular cognitive architecture** called LLaMAR which integrates planning, acting, correcting and self-verification through distinct LLM roles

Key Contributions:

- A **modular cognitive architecture** called LLaMAR which integrates planning, acting, correcting and self-verification through distinct LLM roles
- MAP-THOR a **test-suite and a benchmark** on language-based multi-agent robotic planning based on AI2THOR

Key Contributions:

- A **modular cognitive architecture** called LLaMAR which integrates planning, acting, correcting and self-verification through distinct LLM roles
- MAP-THOR a **test-suite and a benchmark** on language-based multi-agent robotic planning based on AI2THOR
- LLaMAR can create **performant long-horizon planning** in multi-agent tasks by creating subtasks and assigning them to different agents